

m7vxdqs4n

April 18, 2025

## 0.1 Linear regression by using Deep Neural network: Implement Boston housing price prediction problem by Linear regression using Deep Neural Network. Use Boston House Price prediction Dataset.

```
[1]: import tensorflow as tf
      from tensorflow import keras
      from tensorflow.keras import layers
      import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      from sklearn.preprocessing import StandardScaler
      from sklearn.model_selection import train_test_split
      from tensorflow.keras.optimizers import Adam
```

```
[2]: df = pd.read_csv("HousingData.csv")
      df.head()
```

```
[2]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	NaN	36.2

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
---
```

```

0    CRIM      486 non-null    float64
1     ZN      486 non-null    float64
2    INDUS    486 non-null    float64
3     CHAS    486 non-null    float64
4     NOX     506 non-null    float64
5      RM     506 non-null    float64
6     AGE     486 non-null    float64
7     DIS     506 non-null    float64
8     RAD     506 non-null    int64
9     TAX     506 non-null    int64
10    PTRATIO  506 non-null    float64
11     B       506 non-null    float64
12    LSTAT    486 non-null    float64
13    MEDV     506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB

```

```
[4]: df.fillna(df.mean(), inplace=True)
```

```
[5]: X = df.drop(columns=['MEDV'])
     y = df['MEDV']
```

```
[6]: scaler = StandardScaler()
     X = scaler.fit_transform(X)
```

```
[7]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
     ↪random_state=42)
```

```
[8]: model = keras.Sequential([
     keras.Input(shape=(X.shape[1],)), # Explicit Input Layer
     layers.Dense(64, activation='relu'), # Hidden Layer 1
     layers.Dense(32, activation='relu'), # Hidden Layer 2
     layers.Dense(1, activation='linear') # Output layer (Regression)
])
```

```
[9]: model.compile(optimizer=Adam(learning_rate=0.01), loss='mse', metrics=['mae'])
```

```
[10]: history = model.fit(X_train, y_train, epochs=100, validation_data=(X_test,
     ↪y_test), batch_size=16, verbose=1)
```

```

Epoch 1/100
26/26          1s 8ms/step - loss:
440.8634 - mae: 18.5338 - val_loss: 57.2412 - val_mae: 5.5565
Epoch 2/100
26/26          0s 2ms/step - loss:
45.3372 - mae: 5.1544 - val_loss: 21.6366 - val_mae: 3.2939
Epoch 3/100
26/26          0s 3ms/step - loss:

```

```
5.3956 - mae: 1.7245 - val_loss: 11.9814 - val_mae: 2.3072
Epoch 100/100
26/26          0s 2ms/step - loss:
3.7539 - mae: 1.5173 - val_loss: 11.3920 - val_mae: 2.2617
```


```
[11]: loss, mae = model.evaluate(X_test, y_test)
      print(f"Test Loss (MSE): {loss}")
      print(f"Test Mean Absolute Error (MAE): {mae}")
```

```
4/4          0s 0s/step - loss:
8.9360 - mae: 2.1187
Test Loss (MSE): 11.391955375671387
Test Mean Absolute Error (MAE): 2.261707305908203
```

```
[12]: predictions = model.predict(X_test)
```

```
4/4          0s 6ms/step
```

```
[13]: plt.scatter(y_test, predictions)
      plt.xlabel("Actual Prices ($1000s)")
      plt.ylabel("Predicted Prices ($1000s)")
      plt.title("Actual vs Predicted House Prices")
      plt.show()
```



wxs98airg

April 18, 2025

## 0.1 Binary classification using Deep Neural Networks Example: Classify movie reviews into “positive” reviews and “negative” reviews, just based on the text content of the reviews. Use IMDB dataset.

```
[1]: import tensorflow as tf
      from tensorflow import keras
      from tensorflow.keras import layers
      import matplotlib.pyplot as plt
```

```
[2]: vocab_size = 10000
      max_length = 200
      (x_train, y_train), (x_test, y_test) = keras.datasets.imdb.
      ↪load_data(num_words=vocab_size)
```

```
[3]: x_train = keras.preprocessing.sequence.pad_sequences(x_train,
      ↪maxlen=max_length, padding='post')
      x_test = keras.preprocessing.sequence.pad_sequences(x_test, maxlen=max_length,
      ↪padding='post')
```

```
[4]: model = keras.Sequential([
      layers.Embedding(input_dim=vocab_size, output_dim=64), # Removed
      ↪input_length
      layers.Conv1D(32, 5, activation='relu'), # 1D Convolution for feature
      ↪extraction
      layers.GlobalMaxPooling1D(), # Reduce dimensions
      layers.Dense(64, activation='relu'), # Fully connected layer
      layers.Dense(1, activation='sigmoid') # Output layer (Binary
      ↪classification)
  ])
```

```
[5]: model.compile(optimizer='adam', loss='binary_crossentropy',
      ↪metrics=['accuracy'])
```

```
[6]: history = model.fit(x_train, y_train, epochs=5, validation_data=(x_test,
      ↪y_test), batch_size=64, verbose=1)
```

Epoch 1/5

391/391


8s 18ms/step -

```
accuracy: 0.6687 - loss: 0.5754 - val_accuracy: 0.8692 - val_loss: 0.3026
Epoch 2/5
391/391          7s 18ms/step -
accuracy: 0.9185 - loss: 0.2122 - val_accuracy: 0.8845 - val_loss: 0.2767
Epoch 3/5
391/391          7s 18ms/step -
accuracy: 0.9827 - loss: 0.0740 - val_accuracy: 0.8825 - val_loss: 0.3220
Epoch 4/5
391/391          7s 18ms/step -
accuracy: 0.9979 - loss: 0.0180 - val_accuracy: 0.8816 - val_loss: 0.3672
Epoch 5/5
391/391          7s 17ms/step -
accuracy: 1.0000 - loss: 0.0031 - val_accuracy: 0.8836 - val_loss: 0.4037
```

```
[7]: loss, accuracy = model.evaluate(x_test, y_test)
     print(f"Test Accuracy: {accuracy:.4f}")
     print(f"Test Loss: {loss:.4f}")
```

```
782/782          3s 4ms/step -
accuracy: 0.8828 - loss: 0.4049
Test Accuracy: 0.8836
Test Loss: 0.4037
```

```
[8]: plt.plot(history.history['accuracy'], label='Train Accuracy')
     plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
     plt.xlabel('Epochs')
     plt.ylabel('Accuracy')
     plt.legend()
     plt.title('Training and Validation Accuracy')
     plt.show()
```



rbezg8v4z

April 18, 2025

## 0.1 Convolutional neural network (CNN) (Any One from the following)

##### Use any dataset of plant disease and design a plant disease detection system using CNN.  
##### Use MNIST Fashion Dataset and create a classifier to classify fashion clothing into categories.

```
[1]: import pandas as pd
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, models
```

```
[2]: train_df = pd.read_csv('fashion-mnist_train.csv')
test_df = pd.read_csv('fashion-mnist_test.csv')
```

```
[3]: train_df.head(2)
```

```
[3]:   label  pixel1  pixel2  pixel3  pixel4  pixel5  pixel6  pixel7  pixel8  \
0      2      0      0      0      0      0      0      0      0
1      9      0      0      0      0      0      0      0      0

      pixel9  ...  pixel775  pixel776  pixel777  pixel778  pixel779  pixel780  \
0      0  ...      0      0      0      0      0      0
1      0  ...      0      0      0      0      0      0

      pixel781  pixel782  pixel783  pixel784
0      0      0      0      0
1      0      0      0      0

[2 rows x 785 columns]
```

```
[4]: test_df.head(2)
```

```
[4]:   label  pixel1  pixel2  pixel3  pixel4  pixel5  pixel6  pixel7  pixel8  \
0      0      0      0      0      0      0      0      0      9
1      1      0      0      0      0      0      0      0      0

      pixel9  ...  pixel775  pixel776  pixel777  pixel778  pixel779  pixel780  \
0      8  ...     103      87      56      0      0      0
```

```
1      0 ...      34      0      0      0      0      0
```

```
      pixel781 pixel782 pixel783 pixel784
0          0          0          0          0
1          0          0          0          0
```

```
[2 rows x 785 columns]
```

```
[5]: # Split features and labels
x_train = train_df.iloc[:, 1:].values
y_train = train_df.iloc[:, 0].values
```

```
[6]: x_test = test_df.iloc[:, 1:].values
y_test = test_df.iloc[:, 0].values
```

```
[7]: # Normalize pixel values
x_train = x_train / 255.0
x_test = x_test / 255.0
```

```
[8]: # Reshape for CNN: (samples, height, width, channels)
x_train = x_train.reshape(-1, 28, 28, 1)
x_test = x_test.reshape(-1, 28, 28, 1)
```

```
[9]: # Build CNN model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D(2, 2),

    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

C:\Python312\Lib\site-packages\keras\src\layers\convolutional\base\_conv.py:107:  
UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When  
using Sequential models, prefer using an `Input(shape)` object as the first  
layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
[10]: # Compile model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```



```
[11]: # Train
model.fit(x_train, y_train, epochs=5, validation_split=0.1)
```

```
Epoch 1/5
1688/1688          13s 7ms/step -
accuracy: 0.7610 - loss: 0.6718 - val_accuracy: 0.8638 - val_loss: 0.3871
Epoch 2/5
1688/1688          12s 7ms/step -
accuracy: 0.8786 - loss: 0.3398 - val_accuracy: 0.8882 - val_loss: 0.3101
Epoch 3/5
1688/1688          13s 7ms/step -
accuracy: 0.8976 - loss: 0.2819 - val_accuracy: 0.8890 - val_loss: 0.3041
Epoch 4/5
1688/1688          13s 8ms/step -
accuracy: 0.9079 - loss: 0.2526 - val_accuracy: 0.8978 - val_loss: 0.2800
Epoch 5/5
1688/1688          12s 7ms/step -
accuracy: 0.9167 - loss: 0.2238 - val_accuracy: 0.8975 - val_loss: 0.2802
```

```
[11]: <keras.src.callbacks.history.History at 0x25f324b7050>
```

```
[12]: # Evaluate
loss, acc = model.evaluate(x_test, y_test)
print(f"\nTest Accuracy: {acc}")
```

```
313/313          1s 4ms/step -
accuracy: 0.9019 - loss: 0.2680
```

```
Test Accuracy: 0.9064000248908997
```

```
[13]: import matplotlib.pyplot as plt
class_names=['T-shirt/
↳top', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle_
↳boot']
plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(x_train[i], cmap=plt.cm.binary)
    plt.xlabel(class_names[y_train[i]])
plt.show()
```



[ ]:

//Design and implement Parallel Breadth First Search and Depth First Search based on existing algorithms using OpenMP. Use a Tree or an undirected graph for BFS and DFS .

```
#include <iostream>
#include <vector>
#include <queue>
#include <stack>
#include <omp.h>
using namespace std;

const int N = 6; // Number of nodes
vector<int> graph[N];
bool visited_bfs[N], visited_dfs[N];

// Add edge to undirected graph
void addEdge(int u, int v) {
    graph[u].push_back(v);
    graph[v].push_back(u);
}

// Parallel BFS using OpenMP
void parallelBFS(int start) {
    queue<int> q;
    q.push(start);
    visited_bfs[start] = true;

    while (!q.empty()) {
        int size = q.size();

        #pragma omp parallel for
        for (int i = 0; i < size; i++) {
            int node;
            #pragma omp critical
            {
                node = q.front(); q.pop();
                cout << "BFS visited: " << node << endl;
            }

            for (int neighbor : graph[node]) {
                #pragma omp critical
                {
                    if (!visited_bfs[neighbor]) {
                        visited_bfs[neighbor] = true;
                        q.push(neighbor);
                    }
                }
            }
        }
    }
}

// Parallel DFS using OpenMP
void parallelDFS(int start) {
```

```

stack<int> s;
s.push(start);
visited_dfs[start] = true;

while (!s.empty()) {
    int node;
    #pragma omp critical
    {
        node = s.top(); s.pop();
        cout << "DFS visited: " << node << endl;
    }

    #pragma omp parallel for
    for (int i = 0; i < graph[node].size(); i++) {
        int neighbor = graph[node][i];
        #pragma omp critical
        {
            if (!visited_dfs[neighbor]) {
                visited_dfs[neighbor] = true;
                s.push(neighbor);
            }
        }
    }
}

int main() {
    addEdge(0, 1);
    addEdge(0, 2);
    addEdge(1, 3);
    addEdge(1, 4);
    addEdge(2, 5);

    cout << "Parallel BFS:\n";
    parallelBFS(0);

    cout << "\nParallel DFS:\n";
    parallelDFS(0);

    return 0;
}

//run= g++ -fopenmp HPC_Practical_1.cpp -o HPC_Practical_1
//Windows:- HPC_Practical_1.exe
//Linux:- ./HPC_Practical_1

```

//Write a program to implement Parallel Bubble Sort and Merge sort using OpenMP.  
Use existing algorithms and measure the performance of sequential and parallel algorithms.

```
#include <iostream>
#include <vector>
#include <omp.h>
using namespace std;

// Sequential Bubble Sort
void bubbleSortSeq(vector<int>& arr) {
    int n = arr.size();
    for (int i = 0; i < n-1; i++)
        for (int j = 0; j < n-i-1; j++)
            if (arr[j] > arr[j+1])
                swap(arr[j], arr[j+1]);
}

// Parallel Bubble Sort
void bubbleSortPar(vector<int>& arr) {
    int n = arr.size();
    for (int i = 0; i < n; i++) {
        #pragma omp parallel for
        for (int j = i % 2; j < n - 1; j += 2)
            if (arr[j] > arr[j + 1])
                swap(arr[j], arr[j + 1]);
    }
}

// Merge function
void merge(vector<int>& arr, int l, int m, int r) {
    vector<int> left(arr.begin() + l, arr.begin() + m + 1);
    vector<int> right(arr.begin() + m + 1, arr.begin() + r + 1);
    int i = 0, j = 0, k = l;
    while (i < left.size() && j < right.size())
        arr[k++] = (left[i] < right[j]) ? left[i++] : right[j++];
    while (i < left.size()) arr[k++] = left[i++];
    while (j < right.size()) arr[k++] = right[j++];
}

// Sequential Merge Sort
void mergeSortSeq(vector<int>& arr, int l, int r) {
    if (l < r) {
        int m = (l + r) / 2;
        mergeSortSeq(arr, l, m);
        mergeSortSeq(arr, m + 1, r);
        merge(arr, l, m, r);
    }
}

// Parallel Merge Sort
void mergeSortPar(vector<int>& arr, int l, int r) {
    if (l < r) {
        int m = (l + r) / 2;
```

```

        #pragma omp parallel sections
        {
            #pragma omp section
            mergeSortPar(arr, l, m);
            #pragma omp section
            mergeSortPar(arr, m + 1, r);
        }
        merge(arr, l, m, r);
    }
}

int main() {
    vector<int> data = {8, 5, 2, 9, 1, 4};
    vector<int> arr1 = data, arr2 = data;
    vector<int> arr3 = data, arr4 = data;

    bubbleSortSeq(arr1);
    bubbleSortPar(arr2);
    mergeSortSeq(arr3, 0, arr3.size() - 1);
    mergeSortPar(arr4, 0, arr4.size() - 1);

    cout << "Sorted (Seq Bubble): ";
    for (int x : arr1) cout << x << " ";
    cout << "\nSorted (Par Bubble): ";
    for (int x : arr2) cout << x << " ";
    cout << "\nSorted (Seq Merge): ";
    for (int x : arr3) cout << x << " ";
    cout << "\nSorted (Par Merge): ";
    for (int x : arr4) cout << x << " ";
}

```

//Compile Windows:- g++ -fopenmp HPC\_Practical\_2.cpp -o HPC\_Practical\_2

//Run:- HPC\_Practical\_2.exe

//Compile linux:- sudo apt update

//sudo apt install g++ libomp-dev

//g++ -fopenmp HPC\_Practical\_2.cpp -o HPC\_Practical\_2

///  
//./HPC\_Practical\_2

bqcofxpye

April 18, 2025

## 0.1 Implement Min, Max, Sum and Average operations using Parallel Reduction.

```
[1]: import multiprocessing
import random

def parallel_reduction(operation, arr):
    with multiprocessing.Pool() as pool:
        if operation == "min":
            return min(pool.map(min, arr))
        elif operation == "max":
            return max(pool.map(max, arr))
        elif operation == "sum":
            return sum(pool.map(sum, arr))
        elif operation == "avg":
            return sum(pool.map(sum, arr)) / len(arr)

if __name__ == "__main__":
    arr = [random.randint(0, 10000) for _ in range(10000)]
    chunked_arr = [arr[i::multiprocessing.cpu_count()] for i in
↳range(multiprocessing.cpu_count())]

    print(f"Min: {parallel_reduction('min', chunked_arr)}")
    print(f"Max: {parallel_reduction('max', chunked_arr)}")
    print(f"Sum: {parallel_reduction('sum', chunked_arr)}")
    print(f"Average: {parallel_reduction('avg', chunked_arr)}")
```

```
Min: 0
Max: 10000
Sum: 49891056
Average: 6236382.0
```

```
[ ]:
```

**Write a cuda program for**

**1. Addition of two large vectors**

**2. Matrices Multiplication using CUDA C**

from numba import cuda

import numpy as np

@cuda.jit

def add\_vectors(a, b, c):

i = cuda.grid(1)

if i < a.size:

c[i] = a[i] + b[i]

# Host code

N = 1000000

a = np.arange(N, dtype=np.float32)

b = np.arange(N, dtype=np.float32)

c = np.zeros(N, dtype=np.float32)

# Copy to device

d\_a = cuda.to\_device(a)

d\_b = cuda.to\_device(b)

d\_c = cuda.device\_array\_like(c)

# Launch kernel

threadsperblock = 256

blockspergrid = (N + (threadsperblock - 1)) // threadsperblock

add\_vectors[blockspergrid, threadsperblock](d\_a, d\_b, d\_c)

# Copy back

c = d\_c.copy\_to\_host()

print(c[:5])



```
from numba import cuda
import numpy as np
```

```
@cuda.jit
def matmul(A, B, C):
    row, col = cuda.grid(2)
    if row < C.shape[0] and col < C.shape[1]:
        tmp = 0
        for k in range(A.shape[1]):
            tmp += A[row, k] * B[k, col]
        C[row, col] = tmp
```

```
N = 512
A = np.ones((N, N), dtype=np.float32)
B = np.ones((N, N), dtype=np.float32)
C = np.zeros((N, N), dtype=np.float32)
```

```
d_A = cuda.to_device(A)
d_B = cuda.to_device(B)
d_C = cuda.device_array_like(C)
```

```
threadsperblock = (16, 16)
blockspergrid_x = (A.shape[0] + threadsperblock[0] - 1) // threadsperblock[0]
blockspergrid_y = (B.shape[1] + threadsperblock[1] - 1) // threadsperblock[1]
matmul[(blockspergrid_x, blockspergrid_y), threadsperblock](d_A, d_B, d_C)
```

```
C = d_C.copy_to_host()
```

```
print(C[0, 0])
```